

SVM (ForestFire)

classify the Size_Categorie using SVM

month month of the year: 'jan' to 'dec'

day day of the week: 'mon' to 'sun'

FFMC FFMC index from the FWI system: 18.7 to 96.20

DMC DMC index from the FWI system: 1.1 to 291.3

DC DC index from the FWI system: 7.9 to 860.6

ISI ISI index from the FWI system: 0.0 to 56.10

temp temperature in Celsius degrees: 2.2 to 33.30

RH relative humidity in %: 15.0 to 100

wind wind speed in km/h: 0.40 to 9.40

rain outside rain in mm/m2 : 0.0 to 6.4

Size_Categorie the burned area of the forest (Small , Large)

1. Import Libs

In [55]:

```
import numpy as np
import pandas as pd
from sklearn import preprocessing
from sklearn import metrics
import seaborn as sns
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from matplotlib import pyplot as plt
from sklearn.decomposition import PCA
from mlxtend.plotting import plot_decision_regions
from sklearn.metrics import confusion_matrix as cm
from sklearn.metrics import accuracy_score as ac
from sklearn.metrics import classification_report as report, roc_curve
from sklearn.metrics import roc_auc_score
```

2. Import Data


In [2]:

```
forest_fire = pd.read_csv('forestfires.csv')
forest_fire
```

Out[2]:

	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	...	monthfeb	monthjan	r
0	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	...	0	0	
1	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	...	0	0	
2	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	...	0	0	
3	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	...	0	0	
4	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	...	0	0	
...	
512	aug	sun	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	...	0	0	
513	aug	sun	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	...	0	0	
514	aug	sun	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	...	0	0	
515	aug	sat	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	...	0	0	
516	nov	tue	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	...	0	0	

517 rows × 31 columns



3. EDA

In [3]:

```
forest_fire.isna().sum()
```

Out[3]:

month	0
day	0
FFMC	0
DMC	0
DC	0
ISI	0
temp	0
RH	0
wind	0
rain	0
area	0
dayfri	0
daymon	0
daysat	0
daysun	0
daythu	0
daytue	0
daywed	0
monthapr	0
monthaug	0
monthdec	0
monthfeb	0
monthjan	0
monthjul	0
monthjun	0
monthmar	0
monthmay	0
monthnov	0
monthoct	0
monthsep	0
size_category	0

dtype: int64

In [4]:

```
forest_fire.dtypes
```

Out[4]:

month	object
day	object
FFMC	float64
DMC	float64
DC	float64
ISI	float64
temp	float64
RH	int64
wind	float64
rain	float64
area	float64
dayfri	int64
daymon	int64
daysat	int64
daysun	int64
daythu	int64
daytue	int64
daywed	int64
monthapr	int64
monthaug	int64
monthdec	int64
monthfeb	int64
monthjan	int64
monthjul	int64
monthjun	int64
monthmar	int64
monthmay	int64
monthnov	int64
monthoct	int64
monthsep	int64
size_category	object
dtype:	object

In [5]:

```
forest_fire['size_category'].value_counts()
```

Out[5]:

small	378
large	139

Name: size_category, dtype: int64

In [6]:

forest_fire.describe().T

Out[6]:

	count	mean	std	min	25%	50%	75%	max
FFMC	517.0	90.644681	5.520111	18.7	90.2	91.60	92.90	96.20
DMC	517.0	110.872340	64.046482	1.1	68.6	108.30	142.40	291.30
DC	517.0	547.940039	248.066192	7.9	437.7	664.20	713.90	860.60
ISI	517.0	9.021663	4.559477	0.0	6.5	8.40	10.80	56.10
temp	517.0	18.889168	5.806625	2.2	15.5	19.30	22.80	33.30
RH	517.0	44.288201	16.317469	15.0	33.0	42.00	53.00	100.00
wind	517.0	4.017602	1.791653	0.4	2.7	4.00	4.90	9.40
rain	517.0	0.021663	0.295959	0.0	0.0	0.00	0.00	6.40
area	517.0	12.847292	63.655818	0.0	0.0	0.52	6.57	1090.84
dayfri	517.0	0.164410	0.371006	0.0	0.0	0.00	0.00	1.00
daymon	517.0	0.143133	0.350548	0.0	0.0	0.00	0.00	1.00
daysat	517.0	0.162476	0.369244	0.0	0.0	0.00	0.00	1.00
daysun	517.0	0.183752	0.387657	0.0	0.0	0.00	0.00	1.00
daythu	517.0	0.117988	0.322907	0.0	0.0	0.00	0.00	1.00
daytue	517.0	0.123791	0.329662	0.0	0.0	0.00	0.00	1.00
daywed	517.0	0.104449	0.306138	0.0	0.0	0.00	0.00	1.00
monthapr	517.0	0.017408	0.130913	0.0	0.0	0.00	0.00	1.00
monthaug	517.0	0.355899	0.479249	0.0	0.0	0.00	1.00	1.00
monthdec	517.0	0.017408	0.130913	0.0	0.0	0.00	0.00	1.00
monthfeb	517.0	0.038685	0.193029	0.0	0.0	0.00	0.00	1.00
monthjan	517.0	0.003868	0.062137	0.0	0.0	0.00	0.00	1.00
monthjul	517.0	0.061896	0.241199	0.0	0.0	0.00	0.00	1.00
monthjun	517.0	0.032882	0.178500	0.0	0.0	0.00	0.00	1.00
monthmar	517.0	0.104449	0.306138	0.0	0.0	0.00	0.00	1.00
monthmay	517.0	0.003868	0.062137	0.0	0.0	0.00	0.00	1.00
monthnov	517.0	0.001934	0.043980	0.0	0.0	0.00	0.00	1.00
monthoct	517.0	0.029014	0.168007	0.0	0.0	0.00	0.00	1.00
monthsep	517.0	0.332689	0.471632	0.0	0.0	0.00	1.00	1.00

Checking Outlires


In [7]:

```
sns.boxplot(forest_fire['area'])  
plt.grid()
```

C:\Users\shubham\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(  

```




There are 3 Outlier instances in our data

In [8]:

```
plt.figure(figsize=(16,5))  
print("Skewness =",forest_fire['area'].skew())  
print("Kurtosis =",forest_fire['area'].kurtosis())  
sns.kdeplot(forest_fire['area'],shade=True,color='r')  
plt.xticks([i for i in range(0,1200,50)])  
plt.show()
```

Skewness = 12.846933533934868

Kurtosis = 194.1407210942299



The Data is highly skewed and has large kurtosis value

Majority of the forest fires do not cover a large area, most of the damaged area is under 100 hectares of land

In [9]:

```
dfa = forest_fire[forest_fire.columns[0:10]]
month_column = dfa.select_dtypes(include='object')
month_column
```


Out[9]:

	month	day
0	mar	fri
1	oct	tue
2	oct	sat
3	mar	fri
4	mar	sun
...
512	aug	sun
513	aug	sun
514	aug	sun
515	aug	sat
516	nov	tue

517 rows × 2 columns

In [10]:

```
plt.figure(figsize=(16,10))
for i,col in enumerate(month_column,1):
    plt.subplot(2,2,i)
    forest_fire[col].value_counts(normalize=True).plot.bar()
    plt.ylabel(col)
    plt.xlabel('% distribution per category')
plt.tight_layout()
plt.show()
```



In [16]:


```
scl =forest_fire['size_category'] == 'large'
```

In [23]:

```
forest_fire.plot(kind='scatter', x='month', y='day', alpha=0.2, s=20*forest_fire['area'],fi
plt.xlabel('Months',color='red',fontsize=15)
plt.ylabel('Days',color='red',fontsize=15)
plt.title('Different Time periode of fire',color='blue',fontsize=18)
```

Out[23]:

Text(0.5, 1.0, 'Different Time periode of fire')



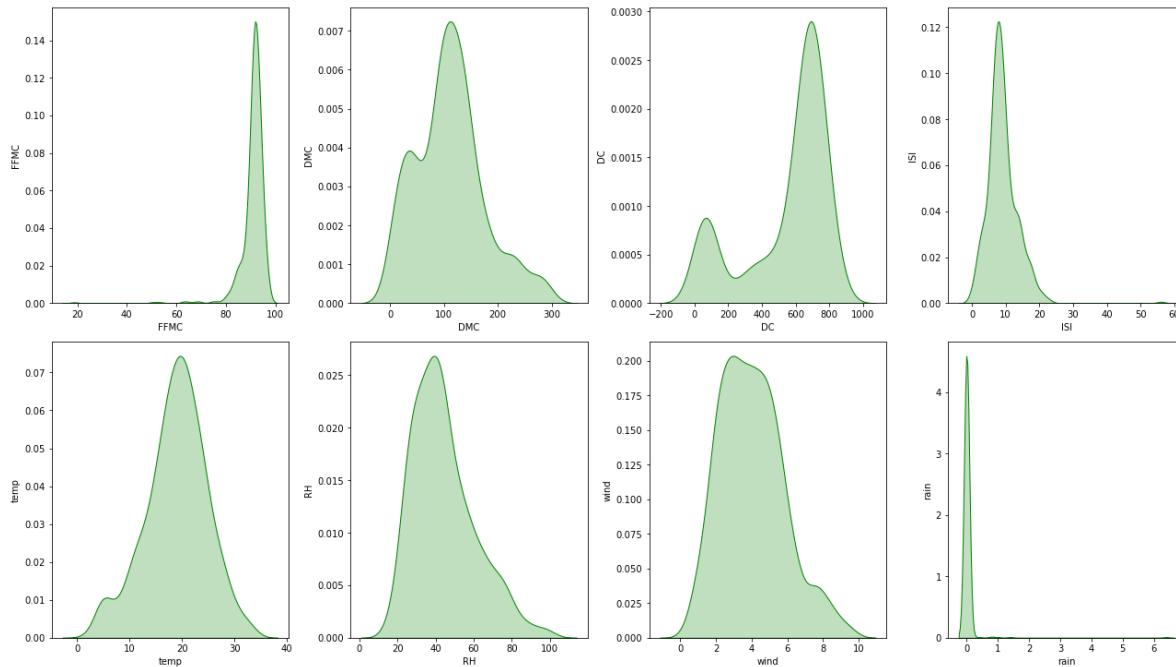
Majority of the fire accors in the month Aug and Sep || sunday and friday have recorded the most cases of fire

In [25]:

```
num_columns = dfa.select_dtypes(exclude='object')
```


In [26]:

```
plt.figure(figsize=(18,40))
for i,col in enumerate(num_columns,1):
    plt.subplot(8,4,i)
    sns.kdeplot(forest_fire[col],color='g',shade=True,legend=True)
    plt.ylabel(col)
plt.tight_layout()
plt.show()
```



In [27]:

```
pd.DataFrame(data=[num_columns.skew(),num_columns.kurtosis()],index=['skewness','kurtosis'])
```

Out[27]:

	FFMC	DMC	DC	ISI	temp	RH	wind	rain
skewness	-6.575606	0.547498	-1.100445	2.536325	-0.331172	0.862904	0.571001	19.816344
kurtosis	67.066041	0.204822	-0.245244	21.458037	0.136166	0.438183	0.054324	421.295964

Finding Correlation

In [28]:


```
corr = forest_fire[forest_fire.columns[0:11]].corr()
```

In [29]:

```
plt.figure(figsize=(10,10))
sns.heatmap(corr,annot=True)
```

Out[29]:

<AxesSubplot:>



4. Model Building

SVM

In [31]:

```
X = forest_fire.iloc[:,2:30]
y = forest_fire.iloc[:,30]
```

In [32]:

X

Out[32]:

	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area	dayfri	...	monthdec	monthfeb
0	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.00	1	...	0	0
1	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.00	0	...	0	0
2	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.00	0	...	0	0
3	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.00	1	...	0	0
4	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.00	0	...	0	0
...
512	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	6.44	0	...	0	0
513	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	54.29	0	...	0	0
514	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	11.16	0	...	0	0
515	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	0.00	0	...	0	0
516	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	0.00	0	...	0	0

517 rows × 28 columns

In [33]:

y

Out[33]:

```
0    small
1    small
2    small
3    small
4    small
...
512  large
513  large
514  large
515  small
516  small
```

Name: size_category, Length: 517, dtype: object

In [34]:

```
mapping = {'small': 0, 'large': 1}
```

In [35]:

```
y = y.replace(mapping)
```

In [36]:

```
x_train,x_test,y_train,y_test = train_test_split(X,y,test_size = 0.20, stratify = y)
```

4.1 Linear

In [37]:


```
model_linear = SVC(kernel = "linear")
model_linear.fit(x_train,y_train)
pred_test_linear = model_linear.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_linear))
```

Accuracy: 0.9711538461538461

In [110]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_linear)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--')
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using Linear Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_linear)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using Linear Kernel for Predicting Size_category



ROC AUC : 0.9690

4.2 Poly

In [38]:


```
model_poly = SVC(kernel = "poly")
model_poly.fit(x_train,y_train)
pred_test_poly = model_poly.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_poly))
```

Accuracy: 0.7692307692307693

In [109]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_poly)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--' )
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using Polynomial Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_poly)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using Polynomial Kernel for Predicting Size_category



ROC AUC : 0.5714

4.3 RBF

In [39]:


```
model_rbf = SVC(kernel = "rbf")
model_rbf.fit(x_train,y_train)
pred_test_rbf = model_rbf.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_rbf))
```

Accuracy: 0.7596153846153846

In [107]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_rbf)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--' )
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using RBF Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_rbf)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using RBF Kernel for Predicting Size_category



ROC AUC : 0.5536

4.4 Sigmoid

In [40]:


```
model_sigmoid = SVC(kernel = "sigmoid")
model_sigmoid.fit(x_train,y_train)
pred_test_sigmoid = model_sigmoid.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_sigmoid))
```

Accuracy: 0.7019230769230769

In [108]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_sigmoid)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--' )
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using Sigmoid Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_sigmoid)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using Sigmoid Kernel for Predicting Size_category



ROC AUC : 0.5141

The Linear Model has best accuracy compare to other Models

In []:

