

# SVM (Salary\_Data)

1) Prepare a classification model using SVM for salary data

Data Description:

age -- age of a person  
workclass -- A work class is a grouping of work  
education -- Education of an individuals  
maritalstatus -- Marital status of an individulas  
occupation -- occupation of an individuals  
relationship --  
race -- Race of an Individual  
sex -- Gender of an Individual  
capitalgain -- profit received from the sale of an investment  
capitalloss -- A decrease in the value of a capital asset  
hoursperweek -- number of hours work per week  
native -- Native of an individual  
Salary -- salary of an individual

## 1. Import Libs

In [1]:

```
import numpy as np
import pandas as pd
from sklearn import preprocessing
from sklearn import metrics
import seaborn as sns
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from matplotlib import pyplot as plt
from sklearn.decomposition import PCA
from mlxtend.plotting import plot_decision_regions
from sklearn.metrics import confusion_matrix as cm
from sklearn.metrics import accuracy_score as ac
from sklearn.metrics import classification_report as report, roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.preprocessing import LabelEncoder
```

## 2. Import Data

In [2]:

```
test_data = pd.read_csv('SalaryData_Test(1).csv')
train_data = pd.read_csv('SalaryData_Train(1).csv')
```

In [3]:

```
test_data
```

Out[3]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White
...	...	...	...	...	...	...	...	...
15055	33	Private	Bachelors	13	Never-married	Prof-specialty	Own-child	White
15056	39	Private	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
15057	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
15058	44	Private	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
15059	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

15060 rows × 14 columns



In [4]:

```
train_data
```

Out[4]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...	...	...	...	...
30156	27	Private	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White
30157	40	Private	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White
30158	58	Private	HS-grad	9	Widowed	Adm-clerical	Unmarried	White
30159	22	Private	HS-grad	9	Never-married	Adm-clerical	Own-child	White
30160	52	Self-emp-inc	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White

30161 rows × 14 columns



### 3. Data Preprocessing

In [5]:

```
df_temp = test_data.append(train_data)
df_temp
```


C:\Users\shubham\AppData\Local\Temp\ipykernel\_2088\1463285300.py:1: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
df_temp = test_data.append(train_data)
```

Out[5]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White
...	...	...	...	...	...	...	...	...
30156	27	Private	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White
30157	40	Private	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White
30158	58	Private	HS-grad	9	Widowed	Adm-clerical	Unmarried	White
30159	22	Private	HS-grad	9	Never-married	Adm-clerical	Own-child	White
30160	52	Self-emp-inc	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White

45221 rows × 14 columns



In [6]:

```
train = train_data.copy()
test = test_data.copy()
test.head()
```

Out[6]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White	Male

In [7]:

```
train.head()
```

Out[7]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female

In [8]:

```
str_c = ["workclass", "education", "maritalstatus", "occupation", "relationship", "race", "sex", "native"]
```

Out[8]:

```
['workclass',
 'education',
 'maritalstatus',
 'occupation',
 'relationship',
 'race',
 'sex',
 'native']
```

In [9]:

```
for i in str_c:
    train[i]= LabelEncoder().fit_transform(train[i])
    test[i]=LabelEncoder().fit_transform(test[i])

mapping = {' >50K': 1, ' <=50K': 0}
train = train.replace({'Salary': mapping})
test = test.replace({'Salary': mapping})
```

In [10]:

```
test.head()
```

Out[10]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	25	2	1	7	4	6	3	2	1
1	38	2	11	9	2	4	0	4	1
2	28	1	7	12	2	10	0	4	1
3	44	2	15	10	2	6	0	2	1
4	34	2	0	6	4	7	1	4	1

In [11]:

```
train.head()
```

Out[11]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	39	5	9	13	4	0	1	4	1
1	50	4	9	13	2	3	0	4	1
2	38	2	11	9	0	5	1	4	1
3	53	2	1	7	2	5	0	2	1
4	28	2	9	13	2	9	5	2	0

In [12]:

```
df = train.append(test)
```

C:\Users\shubham\AppData\Local\Temp\ipykernel\_2088\2767323617.py:1: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
df = train.append(test)
```

In [13]:

```
df1 = df.copy()
df1
```

Out[13]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	s
0	39	5	9	13	4	0	1	4	
1	50	4	9	13	2	3	0	4	
2	38	2	11	9	0	5	1	4	
3	53	2	1	7	2	5	0	2	
4	28	2	9	13	2	9	5	2	
...	...	...	...	...	...	...	...	...	...
15055	33	2	9	13	4	9	3	4	
15056	39	2	9	13	0	9	1	4	
15057	38	2	9	13	2	9	0	4	
15058	44	2	9	13	0	0	3	1	
15059	35	3	9	13	2	3	0	4	

45221 rows × 14 columns

In [14]:

```
df1.describe()
```

Out[14]:

	age	workclass	education	educationno	maritalstatus	occupation	r
count	45221.000000	45221.000000	45221.000000	45221.000000	45221.000000	45221.000000	45
mean	38.548086	2.204507	10.313217	10.118463	2.585148	5.969572	
std	13.217981	0.958132	3.816992	2.552909	1.500460	4.026444	
min	17.000000	0.000000	0.000000	1.000000	0.000000	0.000000	
25%	28.000000	2.000000	9.000000	9.000000	2.000000	2.000000	
50%	37.000000	2.000000	11.000000	10.000000	2.000000	6.000000	
75%	47.000000	2.000000	12.000000	13.000000	4.000000	9.000000	
max	90.000000	6.000000	15.000000	16.000000	6.000000	13.000000	

In [15]:

```
df1.isna().sum()
```

Out[15]:

```
age                0
workclass          0
education          0
educationno        0
maritalstatus      0
occupation         0
relationship       0
race              0
sex               0
capitalgain        0
capitalloss        0
hoursperweek       0
native            0
Salary            0
dtype: int64
```

## Finding Correlation


In [16]:

```
corr = df1.corr()
```



In [17]:

```
plt.figure(figsize=(10,10))
sns.heatmap(corr,annot=True)
plt.show()
```



In [18]:

```
df1.index.is_unique
```

Out[18]:

False

In [19]:


```
df1=df1.loc[~df1.index.duplicated(), :]
```

In [20]:

```
plt.figure(figsize=(16,5))
print("Skew: {}".format(df1['educationno'].skew()))
print("Kurtosis: {}".format(df1['educationno'].kurtosis()))
sns.kdeplot(df1['educationno'],shade=True,color='r')
plt.show()
```

Skew: -0.305378355820322

Kurtosis: 0.643604835875955



**The Data is negatively skewed and has Low Kurtosis value**

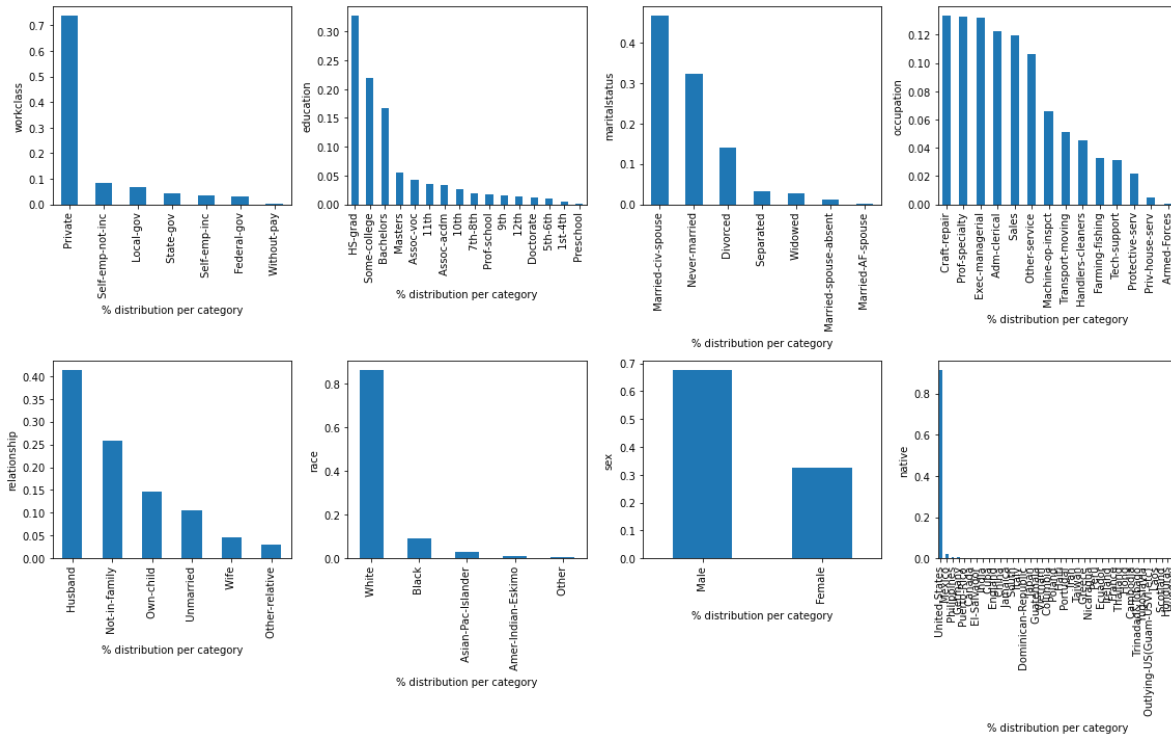
**Most of people have education Number of years of education 8 - 11**

In [21]:

```
dfa = df_temp[df_temp.columns[0:13]]
obj_colum = dfa.select_dtypes(include='object')
```

In [22]:

```
plt.figure(figsize=(16,10))
for i,col in enumerate(obj_colum,1):
    plt.subplot(2,4,i)
    df_temp[col].value_counts(normalize=True).plot.bar()
    plt.ylabel(col)
    plt.xlabel('% distribution per category')
plt.tight_layout()
plt.show()
```

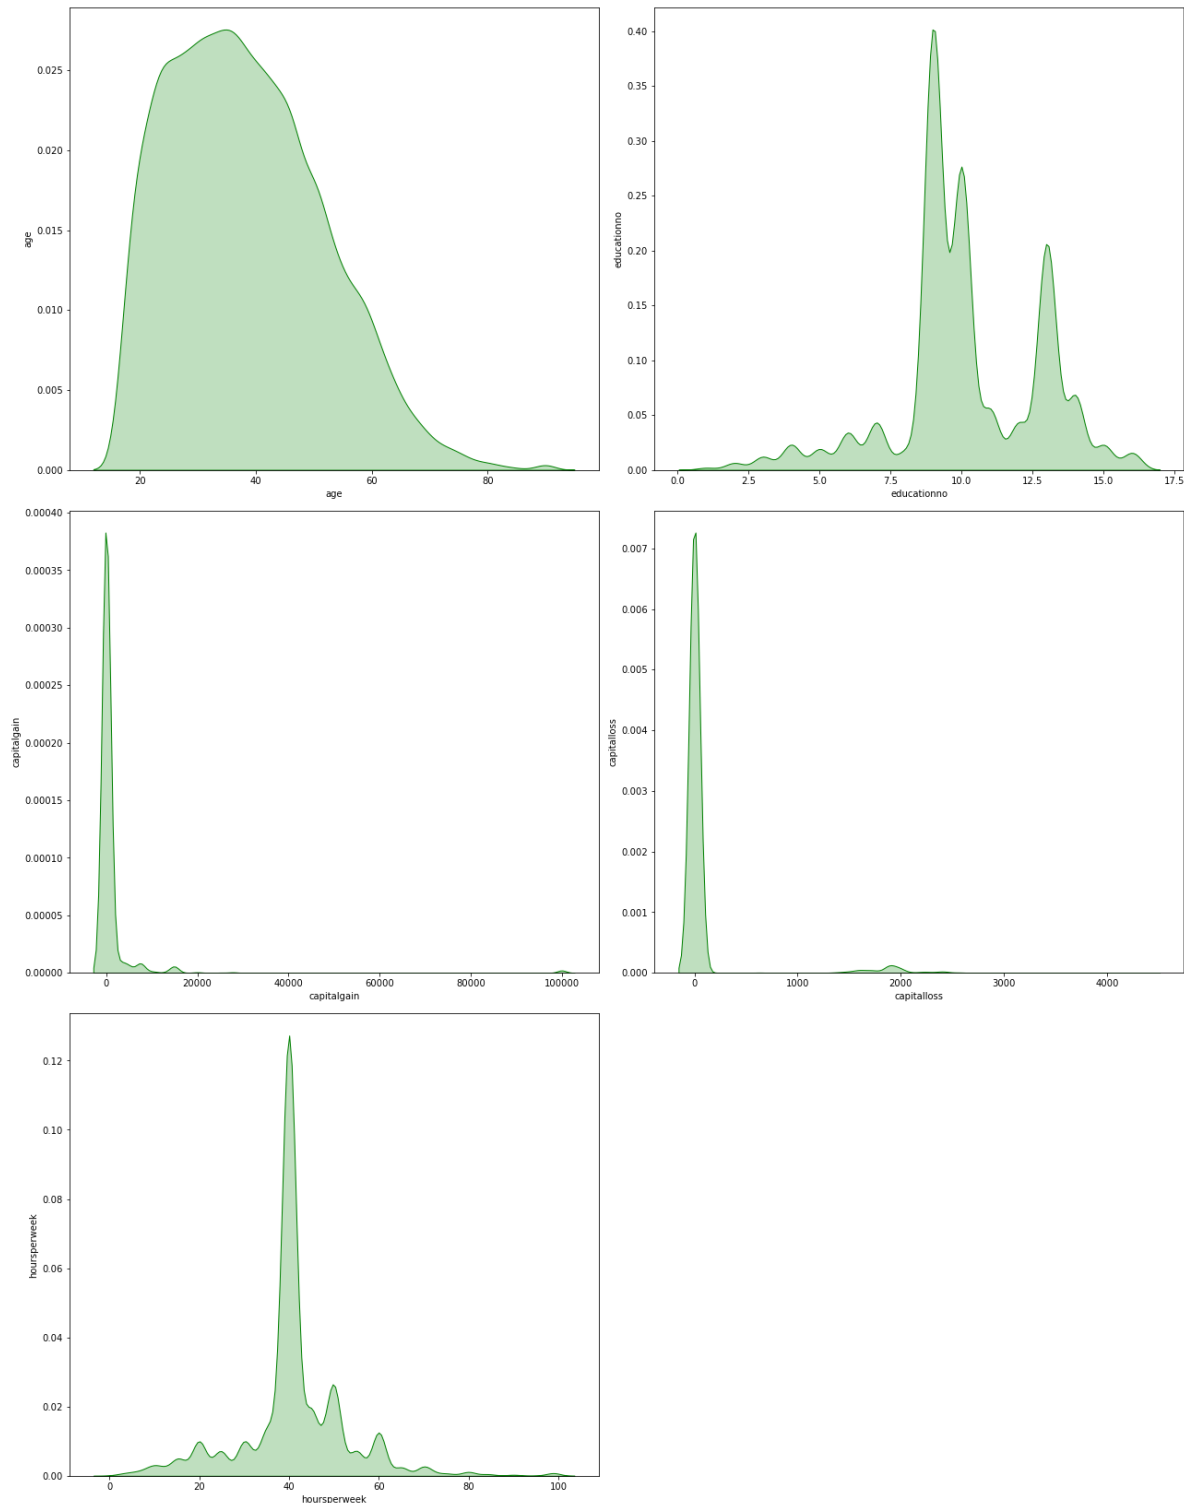


In [23]:

```
num_columns = dfa.select_dtypes(exclude='object')
```

In [24]:

```
plt.figure(figsize=(18,30))
for i,col in enumerate(num_columns,1):
    plt.subplot(4,2,i)
    sns.kdeplot(df1[col],color='g',shade=True,legend=True)
    plt.ylabel(col)
plt.tight_layout()
plt.show()
```



In [25]:

```
pd.DataFrame(data=[num_columns.skew(),num_columns.kurtosis()],index=['skewness','kurtosis'])
```

Out[25]:

	age	educationno	capitalgain	capitalloss	hoursperweek
<b>skewness</b>	0.532784	-0.310621	11.788871	4.517536	0.340536
<b>kurtosis</b>	-0.155931	0.635045	150.147899	19.376085	3.201287

## 4. Model Building

### SVM

In [26]:

```
col = df1.columns
col
```

Out[26]:

```
Index(['age', 'workclass', 'education', 'educationno', 'maritalstatus',
      'occupation', 'relationship', 'race', 'sex', 'capitalgain',
      'capitalloss', 'hoursperweek', 'native', 'Salary'],
      dtype='object')
```

In [69]:

```
x_train = train[col[0:13]]
y_train = train[col[13]]
x_test = test[col[0:13]]
y_test = test[col[13]]
```

In [70]:

```
def norm_func(i):
    x = (i-i.min())/(i.max()-i.min())
    return (x)
```

In [71]:

```
x_train = norm_func(x_train)
x_test = norm_func(x_test)
```

## 4.1 Linear

In [72]:


```
model_linear = SVC(kernel = "linear")
model_linear.fit(x_train,y_train)
pred_test_linear = model_linear.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_linear))
```

Accuracy: 0.8098273572377158

In [73]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_linear)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--')
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using Linear Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_linear)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using Linear Kernel for Predicting Size\_category



ROC AUC : 0.6455

## 4.2 Poly

In [74]:


```
model_poly = SVC(kernel = "poly")
model_poly.fit(x_train,y_train)
pred_test_poly = model_poly.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_poly))
```

Accuracy: 0.8435590969455511

In [75]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_poly)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--' )
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using Polynomial Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_poly)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using Polynomial Kernel for Predicting Size\_category



ROC AUC : 0.7485

## 4.3 RBF

In [76]:


```
model_rbf = SVC(kernel = "rbf")
model_rbf.fit(x_train,y_train)
pred_test_rbf = model_rbf.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_rbf))
```

Accuracy: 0.8432934926958832

In [77]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_rbf)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--' )
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using RBF Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_rbf)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using RBF Kernel for Predicting Size\_category



ROC AUC : 0.7445

## 4.4 Sigmoid

In [78]:

```
model_sigmoid = SVC(kernel = "sigmoid")
model_sigmoid.fit(x_train,y_train)
pred_test_sigmoid = model_sigmoid.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, pred_test_sigmoid))
```


Accuracy: 0.5768924302788845



In [79]:

```
fpr, tpr, thresholds = roc_curve(y_test, pred_test_sigmoid)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, linewidth=2, color='red')
plt.plot([0,1], [0,1], 'k--')
plt.rcParams['font.size'] = 12
plt.title('ROC curve for SVM Classifier using Sigmoid Kernel for Predicting Size_category')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
ROC_AUC = roc_auc_score(y_test, pred_test_sigmoid)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC curve for SVM Classifier using Sigmoid Kernel for Predicting Size\_category



ROC AUC : 0.4285

**The Poly Model has best accuracy compare to other Models. but RBF model has almost equal Accuracy to Poly Model**

In [ ]:

